

Hello Uday surya

Senior AI Engineer | Generative AI | LLMs | AI Agents | RAG

Email | Phone | LinkedIn | GitHub | Portfolio

PROFESSIONAL SUMMARY

Senior AI Engineer with 6+ years of experience designing and building scalable AI-powered applications, intelligent automation systems, and cloud-native backend platforms. Experienced in developing production-grade Generative AI solutions using Large Language Models (LLMs), Retrieval-Augmented Generation (RAG), AI Agents, and vector databases. Strong expertise in Python, FastAPI, Azure OpenAI, LangChain, LangGraph, PostgreSQL, Docker, Kubernetes, and cloud-native architectures. Passionate about building enterprise AI systems that improve productivity, automate business workflows, and deliver intelligent user experiences.

TECHNICAL SKILLS

Programming Languages

Python • C# • JavaScript • TypeScript • SQL

AI & Large Language Models

OpenAI API • Azure OpenAI • Anthropic Claude • Gemini • LangChain • LangGraph • AI Agents • Prompt Engineering • Function Calling • Structured Outputs • MCP (Model Context Protocol)

Retrieval-Augmented Generation (RAG)

Vector Embeddings • pgvector • Pinecone • ChromaDB • Weaviate • Qdrant • Semantic Search • Hybrid Search • Knowledge Retrieval

Backend

FastAPI • Django • Flask • REST APIs • Microservices • Async Programming

Databases

PostgreSQL • SQL Server • Redis • MongoDB

Cloud & DevOps

Microsoft Azure • Azure AI Search • Azure App Service • Azure Functions • Docker • Kubernetes • GitHub Actions • CI/CD • Terraform

Frontend

React • Next.js • HTML • CSS • Tailwind CSS

Tools

Git • GitHub • Postman • VS Code • Linux • Playwright

PROFESSIONAL EXPERIENCE

Senior AI Engineer

Jan 2023 – Present

- Designed and deployed enterprise-grade Generative AI applications using Azure OpenAI and OpenAI GPT models.
- Built Retrieval-Augmented Generation (RAG) systems using vector databases, semantic search, and enterprise document repositories.
- Developed AI Agents capable of multi-step reasoning, tool usage, document retrieval, and workflow automation.
- Designed scalable FastAPI microservices serving AI inference with asynchronous processing.
- Integrated vector databases including pgvector and Pinecone for semantic similarity search.
- Built AI-powered chat assistants capable of answering enterprise knowledge base questions.
- Implemented prompt engineering strategies and structured output generation to improve response quality.
- Containerized AI services using Docker and deployed applications on Kubernetes clusters.
- Designed CI/CD pipelines using GitHub Actions for automated deployments.
- Improved AI response latency by optimizing embedding pipelines and retrieval strategies.

PROJECTS

AI Career Copilot

- Built an AI-powered career assistant capable of discovering jobs, analyzing resumes, and recommending the best job matches.
- Implemented semantic resume matching using OpenAI Embeddings and PostgreSQL pgvector.
- Integrated Microsoft Careers API for automated AI job discovery.
- Developed intelligent resume analysis using OCR, PDF parsing, and LLM-powered skill extraction.
- Built REST APIs using FastAPI.

- Technologies:

Python, FastAPI, OpenAI, LangChain, PostgreSQL, pgvector, Docker, Azure

Education

MONTH 2014 - MONTH 2018

JNTUH

Bachelor of Technology

Computer Science and Engineering

CERTIFICATIONS

Microsoft Certified: Azure AI Fundamentals (AI-900)

Microsoft Certified: Azure Fundamentals (AZ-900)

Microsoft Certified: Azure AI Engineer Associate (AI-102)

OpenAI API Developer

Docker & Kubernetes

LangChain for LLM Application Development